

B461 Assignment 8

Relational Programming

1. Consider the relation schema **A(x int)** representing a schema for storing a set of integers A .

Using arrays to represent sets, write a PostgreSQL program

superSetsOfSet(X int[])

that returns each subset of A that is a superset of X , i.e., each set Y such that $X \subseteq Y \subseteq A$.

For example, if $X = \{\}$, then **superSetsofSets(X)** should return each element of the powerset of A .

2. Consider the relation schema **Graph(source int, target int)** representing a schema for storing a directed graph G of edges.

A path in G is a sequence of vertices (v_0, v_1, \dots, v_n) such that, for each $i \in [0, n-1]$, (v_i, v_{i+1}) is an edge of G . We say that such a path connects v_0 to v_n via an n -length path.

- (a) Write a PostgreSQL program **connectedByEvenLengthPath()** that returns the pairs of vertices (s, t) of G that are connected via an *even*-length path in G .
- (b) Write a PostgreSQL program **connectedByOddLengthPath()** that returns the pairs of vertices (s, t) of G that are connected via an *odd*-length path in G .

3. Consider the relation schema **Graph(source int, target int)** representing the schema for storing a directed graph G of edges.

Now let G be a directed graph that is acyclic, a graph without cycles.¹

A *topological sort* of a graph is a list (array) of **all** its vertices (v_1, v_2, \dots, v_n) such that for each edge (s, t) in G , vertex s occurs before vertex t in this list.

Write a PostgreSQL program **topologicalSort()** that returns a topological sort of G .

4. Consider the relation schema **document(doc int, words text[])** representing a relation of pairs (d, W) where d is a unique id denoting a document and W denotes the set of words that occur in d .

Let **W** denote the set of all words that occur in the documents and let t be a positive integer denoting a *threshold*.

Let $X \subseteq \mathbf{W}$. We say that X is t -frequent if

$$\text{count}(\{d \mid (d, W) \in \text{document and } X \subseteq W\}) \geq t$$

¹A cycle is a path (v_0, \dots, v_k) where $v_0 = v_k$.

In other words, X is *t-frequent* if there are at least t documents that contain all the words in X .

Write a PostgreSQL program `frequentSets(t int)` that returns each t -frequent set.

In a good solution for this problem, you should use the following rule: if X is not t -frequent then any set Y such that $X \subseteq Y$ is not t -frequent either. In the literature, this is called the *Apriori* rule of the frequent itemset mining problem.

5. Consider the following relational schemas. (You can assume that the domain of each of the attributes in these relations is `int`.)

`partSubpart(pid,sid,quantity)`
`basicPart(pid,weight)`

A tuple (p, s, q) is in `partSubPart` if part s occurs q times as a **direct** subpart of part p . For example, think of a car c that has 4 wheels w and 1 radio r . Then $(c, w, 4)$ and $(c, r, 1)$ would be in `partSubpart`. Furthermore, then think of a wheel w that has 5 bolts b . Then $(w, b, 5)$ would be in `partSubpart`.

A tuple (p, w) is in `basicPart` if basic part p has weight w . A basic part is defined as a part that does not have subparts. In other words, the pid of a basic part does not occur in the pid column of `partSubpart`.

(In the above example, a bolt and a radio would be basic parts, but car and wheel would not be basic parts.)

We define the *aggregated weight* of a part inductively as follows:

- (a) If p is a basic part then its aggregated weight is its weight as given in the `basicPart` relation
- (b) If p is not a basic part, then its aggregated weight is the sum of the aggregated weights of its subparts, each multiplied by the quantity with which these subparts occur in the `partSubpart` relation.

Example tables: The following example is based on a desk lamp with pid 1. Suppose a desk lamp consists of 4 bulbs (with pid 2) and a frame (with pid 3), and a frame consists of a post (with pid 4) and 2 switches (with pid 5). Furthermore, we will assume that the weight of a bulb is 5, that of a post is 50, and that of a switch is 3.

Then the `partSubpart` and `basicPart` relation would be as follows:

partSubPart			Parts	
pid	sid	quantity	pid	weight
1	2	4	2	5
1	3	1	4	50
3	4	1	5	3
3	5	2		

Then the aggregated weight of a lamp is $4 \times 5 + 1 \times (1 \times 50 + 2 \times 3) = 76$.

Write a PostgreSQL function `aggregatedWeight(p integer)` that returns the aggregated weight of a part `p`.

6. Suppose you have a weighted undirected graph G stored in a ternary table with schema

`Graph(source int, target int, weight int)`

A triple (s, t, w) in `Graph` indicates that $Graph$ has an edge (s, t) whose edge weight is w . (In this problem, we will assume that each edge weight is a positive integer.)

Since the graph is undirected, whenever there is an weighted edge (s, t, w) in G , then (t, s, w) is also a weighted edge in the G . Below is an example of a graph G .

Graph G		
source	target	weight
0	1	2
1	0	2
0	4	10
4	0	10
1	3	3
3	1	3
1	4	7
4	1	7
2	3	4
3	2	4
3	4	5
4	3	5
4	2	6
2	4	6

Implement Dijkstra's Algorithm as a PostgreSQL function `Dijkstra(s integer)` to compute the shortest path lengths (i.e., the distances) from some input vertex s in G to all other vertices in G . `Dijkstra(s integer)` should accept an argument s , the source vertex, and outputs a table `Paths` which represents the pairs (t, d) where d is the shortest distance from s to t . To test your procedure, you can use the graph shown above.

When you apply `Dijkstra(0)`, you should obtain the following `Paths` table:

target	distanceToTarget
0	0
1	2
2	9
3	5
4	9

Hint: You can find the details of Dijkstra's Algorithm in the attached pdf document, but you are not required to exactly follow the pseudocode.